Metrics

Total points 6

1. What would be a logloss value for a binary classification task, if we use constant predictor f(x) = 0.5? Round to two decimal places. **1 point**

Enter answer here

2. The best constant predictor for MAE metric is **1 point**

- [ ] 0.5
- [ ] Target mean
- [ ] Target median
- [ ] Target mode
- [ ] Target 50-th percentile

3. The best constant predictor for mean squared error is **1 point**

- [ ] Target mean
- [ ] Average of the target vector
- [ ] $log(y + 1)$, where y is target vector
- [ ] Target variance

4. The best constant prediction for AUC is **1 point**

- [ ] Target median
- [ ] 0.5
- [ ] 1
- [ ] Any constant will lead to the same AUC value
- [ ] Target mean
- [ ] Target mean divided by target variance

5. Suppose the target metric is R-squared. What optimization loss should we use for our models? **1 point**

- [ ] MSE
- [ ] AUC
- [ ] RMSE
- [ ] RMSLE
- [ ] MAE

6. Calculate AUC for these predictions: **1 point**

target	prediction
1	0.39
0	0.52
1	0.91
1	0.85
1	0.49
0	0.02
0	0.44

Round to 2 decimal places.

Enter answer here

Coursera Honor Code Learn more

I, **Chinmay kumar Das**, understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.

Submit | Save draft

Like | Dislike | Report an issue